

Assignments

Niko Amber Assignments!

Assignment 1

Collaborators: Eliza Epstein.

Problem 1

Install the datasets package on the console below using `install.packages("datasets")`. Now load the library.

Answer: I've loaded the library!

Load the USArrests dataset and rename it `dat`. Note that this dataset comes with R, in the package datasets, so there's no need to load data from your computer. Why is it useful to rename the dataset?

```
dat <- USArrests
```

Answer: It is beneficial to rename the data set, so we can replicate analyses without disturbing the original data set. Additionally, it is nice to rename your data set to know exactly it is called

Problem 2

First I am making states lowercase, to be used as variables later.

```
dat$state <- tolower(rownames(USArrests))
```

This dataset has the state names as row names, so we just want to make them into a new variable. We also make them all lower case, because that will help us draw a map later - the map function requires the states to be lower case.

List the variables contained in the dataset `USArrests`.

The variables contained in the dataset 'USArrests' are Murder, Assault, and Rape. Additionally, the data set shows us what percentage of people live in urban areas.

```
USArrests
```

##	Murder	Assault	UrbanPop	Rape	state
## Alabama	13.2	236	58	21.2	alabama
## Alaska	10.0	263	48	44.5	alaska
## Arizona	8.1	294	80	31.0	arizona
## Arkansas	8.8	190	50	19.5	arkansas
## California	9.0	276	91	40.6	california

## Colorado	7.9	204	78 38.7	colorado
## Connecticut	3.3	110	77 11.1	connecticut
## Delaware	5.9	238	72 15.8	delaware
## Florida	15.4	335	80 31.9	florida
## Georgia	17.4	211	60 25.8	georgia
## Hawaii	5.3	46	83 20.2	hawaii
## Idaho	2.6	120	54 14.2	idaho
## Illinois	10.4	249	83 24.0	illinois
## Indiana	7.2	113	65 21.0	indiana
## Iowa	2.2	56	57 11.3	iowa
## Kansas	6.0	115	66 18.0	kansas
## Kentucky	9.7	109	52 16.3	kentucky
## Louisiana	15.4	249	66 22.2	louisiana
## Maine	2.1	83	51 7.8	maine
## Maryland	11.3	300	67 27.8	maryland
## Massachusetts	4.4	149	85 16.3	massachusetts
## Michigan	12.1	255	74 35.1	michigan
## Minnesota	2.7	72	66 14.9	minnesota
## Mississippi	16.1	259	44 17.1	mississippi
## Missouri	9.0	178	70 28.2	missouri
## Montana	6.0	109	53 16.4	montana
## Nebraska	4.3	102	62 16.5	nebraska
## Nevada	12.2	252	81 46.0	nevada
## New Hampshire	2.1	57	56 9.5	new hampshire
## New Jersey	7.4	159	89 18.8	new jersey
## New Mexico	11.4	285	70 32.1	new mexico
## New York	11.1	254	86 26.1	new york
## North Carolina	13.0	337	45 16.1	north carolina
## North Dakota	0.8	45	44 7.3	north dakota
## Ohio	7.3	120	75 21.4	ohio
## Oklahoma	6.6	151	68 20.0	oklahoma
## Oregon	4.9	159	67 29.3	oregon
## Pennsylvania	6.3	106	72 14.9	pennsylvania
## Rhode Island	3.4	174	87 8.3	rhode island
## South Carolina	14.4	279	48 22.5	south carolina
## South Dakota	3.8	86	45 12.8	south dakota
## Tennessee	13.2	188	59 26.9	tennessee
## Texas	12.7	201	80 25.5	texas
## Utah	3.2	120	80 22.9	utah
## Vermont	2.2	48	32 11.2	vermont
## Virginia	8.5	156	63 20.7	virginia
## Washington	4.0	145	73 26.2	washington
## West Virginia	5.7	81	39 9.3	west virginia
## Wisconsin	2.6	53	66 10.8	wisconsin
## Wyoming	6.8	161	60 15.6	wyoming

Problem 3

What type of variable (from the DVB chapter) is Murder?

Answer: Murder is a categorical variable, it is one of many categories that crime falls into. It is not ordinal, there is no ordering of crime.

What R Type of variable is it?

Answer: 'Murder' is a character variable, it contains information that isn't numeric.

Problem 4


What information is contained in this dataset, in general? What do the numbers mean?

Answer: The Data set USArrests contains data about the rate of arrests for murder, rape and assault per 100,000 residents in each US state in 1973. The data set also includes the percent of the population living in urban cities in each state. The collums represent each type or crime (and urban population percentage) The numbers in each row are the rate of arrests per state (per 100,000).

Problem 5

Draw a (histogram) bar graph of Murder with proper labels and title. I used a bar graph instead of a histogram because I feel that is a better way to represent this data.

```
barplot(USArrests$Murder , names.arg = state.name, las=2, xlab = "States", ylab = "Rate of Murder Arrests",
        main = "Arrests for Murder in the United States 1973")
```



Problem 6

Please summarize Murder quantitatively. What are its mean and median? What is the difference between mean and median? What is a quartile, and why do you think R gives you the 1st Qu. and 3rd Qu.?

I used the summary function to gather this information

```
summary(USArrests$Murder)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.800   4.075   7.250   7.788  11.250  17.400
```

The mean for murder is 7.788 meaning this is the average rate of people (per 100,000) murdered in each US state in 1973. The median for murder is 7.250 meaning in the United States in 1973 half of the states had a rate per 100,000 more than 7.250 and half had fewer. Median is the middle of numbers in a data set, while mean is the average of all numbers in said set. If the data is evenly distributed the median will equal the mean.

A quartile is when the data is divided into four equal parts: the 1st, 2nd, 3rd, and 4th quartile. R gives the 1st and 3rd quartile because it represents the middle of the data. 1st quartile is lowest 25% and 3rd quartile is highest 25%.


Problem 7

Repeat the same steps you followed for **Murder**, for the variables **Assault** and **Rape**. Now plot all three histograms together. You can do this by using the command `par(mfrow=c(3,1))` and then plotting each of the three.

Note: I used bar graphs

```
barplot(USArrests$Murder , names.arg = state.name, las=2, xlab = "States", ylab = "Rate of Murder Arrests",
        main = "Arrests for Murder in the United States 1973")
```


Arrests for Murder in the United States 1973



#Bar plot for rape arrests

```
barplot(USArrests$Rape , names.arg = state.name, las=2, xlab = "States", ylab = "Rate of Arrests for Rape",
        main = "Rate of Arrests for Rape in the United States in 1973")
```


Rate of Arrests for Rape in the United States in 1997



#Bar plot for assault arrests

```
barplot(USArrests$Assault , names.arg = state.name, las=2, xlab = "States", ylab = "Rate of Arrests for Assaults in the United States",
        main = "Rate of Arrests for Assaults in the United States")
```

Rate of Arrests for Assaults in the United States

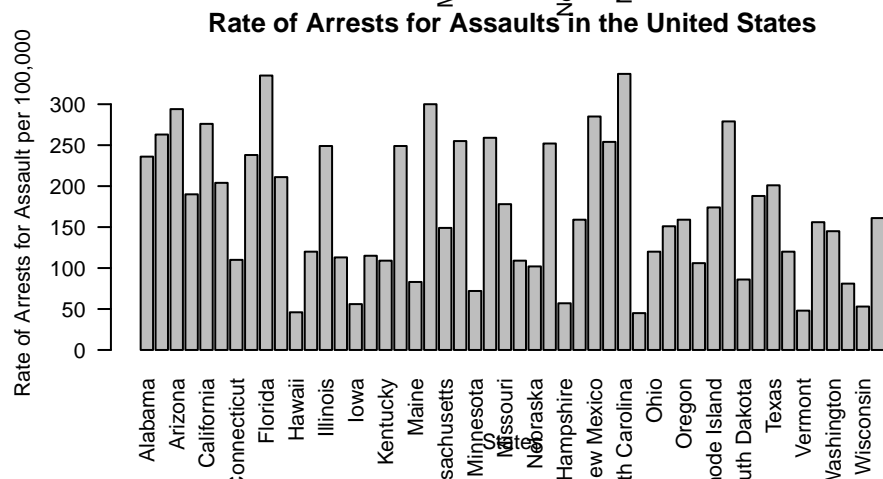




#plotting the graphs together

```
par(mfrow=c(3,1))
barplot(USArrests$Murder , names.arg = state.name, las=2, xlab = "States", ylab = "Rate of Murder Arrests",
        main = "Arrests for Murder in the United States 1973")
```

```
#Bar plot for rape arrests
barplot(USArrests$Rape , names.arg = state.name,las=2, xlab = "States", ylab = "Rate of Arrests for Rape",
        main = "Rate of Arrests for Rape in the United States in 1973")

#Bar plot for assault arrests
barplot(USArrests$Assault , names.arg = state.name,las=2, xlab = "States", ylab = "Rate of Arrests for Assault",
        main = "Rate of Arrests for Assaults in the United States")
```

par

```
## function (... , no.readonly = FALSE)
## {
##     .Pars.readonly <- c("cin", "cra", "csi", "cxy", "din", "page")
```

```
##     single <- FALSE
##     args <- list(...)
##     if (!length(args))
##       args <- as.list(if (no.readonly)
##         .Pars[-match(.Pars.readonly, .Pars)]
##       else .Pars)
##   else {
##     if (all(unlist(lapply(args, is.character))))
##       args <- as.list(unlist(args))
##     if (length(args) == 1) {
##       if (is.list(args[[1L]]) || is.null(args[[1L]]))
##         args <- args[[1L]]
##       else if (is.null(names(args)))
##         single <- TRUE
##     }
##   }
##   value <- .External2(C_par, args)
##   if (single)
##     value <- value[[1L]]
##   if (!is.null(names(args)))
##     invisible(value)
##   else value
## }
## <bytecode: 0x7f84fd681808>
## <environment: namespace:graphics>
```

What does the command `par` do, in your own words (you can look this up by asking R `?par`)?

Answer: This command allows R to set paramaters, this way multiple data sets can be graphed together.

What can you learn from bar graphs the histograms together?

Answer: When we plot these bar graphs together we can compare each state's arrest rates for different crimes. For example, when looking at the bar graphs it is easy to see that North Carolina's arrest rate for assault is much higher than the arrest rate for rape. This can lead researchers to ask questions: why were there more arrests for assaults? Were there possible sexually based assault that should have been charged as rape?

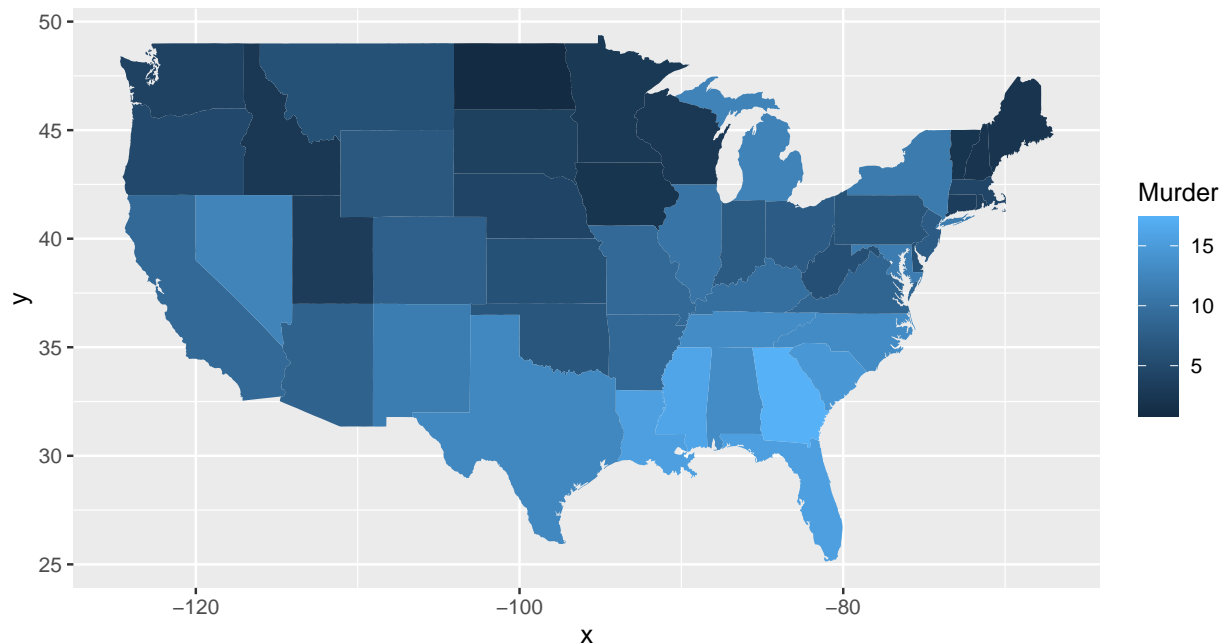
Problem 8

In the console below (not in text), type `install.packages("maps")` and press Enter, and then type `install.packages("ggplot2")` and press Enter. This will install the packages so you can load the libraries.

Run this code:

```
library('maps')
library('ggplot2')

#this code creates a map
ggplot(dat, aes(map_id=state, fill=Murder)) +
  geom_map(map=map_data("state")) +
  expand_limits(x=map_data("state")$long, y=map_data("state")$lat)
```




What does this code do? Explain what each line is doing.

Answer: This code creates a colored map that uses our data set to show murder rates. The lighter the blue, the higher the murder rate.

The first line of code creates a map that is divided by states. The first line also applies the data set to the map, so the color of the state will be lighter if murder rates are higher. The second line of code designs the map by breaking it up by state. The last line of code expands the limits of the map

Here is my histogram

```
hist(dat$Murder, main = "Murder Rate in The United States in 1973", xlab = "Murder Rate", ylab = "Frequency")
```



Assignment 2

Problem 1

```
knitr::opts_chunk$set(echo = TRUE)
```

```
setwd("/Users/nikoamber/Library/Mobile Documents/com~apple~CloudDocs/Penn/Freshman Fall/Criminology/Ass
```

```
dat <- read.csv(file = 'dat.nsduh.small.1.csv')
```

```
names(dat)
```

```
## [1] "mjage"      "cigage"      "iralcage"    "age2"        "sexattract" "speakengl"  
## [7] "irsex"
```

```
nrow(dat)
```

```
## [1] 171
```

```
ncol(dat)
```

```
## [1] 7
```

There are seven columns and 171 rows

Problem 2

Describe the variables in the dataset.

The variables in the data are:

MJAGE which is how old the participants are the first time they used marijuana/hashish. It is a quantitative continuous variable.

CIGAGE which is how old the participants were when they first started smoking cigarettes everyday. It is a quantitative continuous variable.

IRALCAGE which is how old the participants were when they first tried alcohol. it is a quantitative continuous variable.

AGE2S which is the age of the respondents of the survey, categorized into groups. It is a categorical nominal variable.

IRSEX which is which sex participants identify as. It is a categorical, nominal variable. Participants can only identify as male or female. which is coded as 1 or 2.

SEXATTRACT which is who participants of the study are attracted to. It is a categorical, nominal variable. Participants identified their sexual preference on a category scale of 1-6, but there is no ranking.

SPEAKENGL which is how well do participants speak English. It is a categorical, ordinal variable. Participants rate on a scale of 1-4 where 1 is speaking English well and 4 is not at all.


What is this dataset about? Who collected the data, what kind of sample is it, and what was the purpose of generating the data?

This data set is a sample from a 2019 survey taken by the National Survey of Drug Use and Health. This data is only a sample of the survey and has the first 1000 values, and does not include missing values. The data is a random sampling of the US population, and it can be used to tell us more about the US population and their relationship with drugs.

Problem 3: Age and gender

What is the age distribution of the sample like? Make sure you read the codebook to know what the variable values mean.

```
hist(dat$age2, main= "Histogram of ages of participants of 2019 NSDUH survey", xlab= "Ages in Groups", y
```




Problem 3: Age and gender

What is the age distribution of the sample like? Make sure you read the codebook to know what the variable values mean.

```
hist(dat$age2, main= "Histogram of ages of participants of 2019 NSDUH survey", xlab= "Ages in Groups", y
```

Histogram of ages of participants of 2019 NSDUH survey



The age distribution is from ages 12-65+. As mentioned in the code book, it is grouped in sections. but not every group holds the same number of years. Because of this it is difficult to see the age distribution properly. While it appears that most participants are in group 15, group 15 has ages 35-49 (14 years) while groups 1-9 only have one year.

Do you think this age distribution representative of the US population? Why or why not?

Because of the way the ages are grouped together, this histogram is left skewed. However, there are no respondents below the age of 12, which makes sense given the content of this survey. There are also very few older participants. Therefore this data is probably a good representation of the US population for the purposes of this study, but obviously is not representative of the entire US population.

Is the sample balanced in terms of gender? If not, are there more females or males?


```
table(dat$irsex)
```

```
##  
##  1  2  
## 91 80
```

There are more males than females in this sample. It is not representative of the US population, because in the population there are slightly more women than men. This scale is also on a binary, when there are many people in the United States who may not identify with one of these two genders. According to the code book, the entire data set has more females than males, but in these 1000 data points, there are more males than females.

Use this code to draw a stacked bar plot to view the relationship between sex and age. What can you conclude from this plot?

```
tab.agesex <- table(dat$irsex, dat$age2)
barplot(tab.agesex,
        main = "Stacked barchart",
        xlab = "Age category", ylab = "Frequency",
        legend.text = rownames(tab.agesex),
        beside = FALSE) # Stacked bars (default)
```




In the age category 15, that represents respondents between 35 and 49 and 26-29, there is about an even distribution between men and women. However, the younger groups tend to have more female participants than male (group 8 does not have any male respondents) and the groups above 15 seem to have a larger percentage of men.

Problem 4: Substance use

For which of the three substances included in the dataset (marijuana, alcohol, and cigarettes) do individuals tend to use the substance earlier?

```
par(mfrow=c (1,1))
plot(dat$age2, dat$mjage)
```


```
cor(dat$age2, dat$mjage)
```

```
## [1] 0.1811713
```

```
cor(dat$age2, dat$cigage)
```

```
## [1] 0.1697763
```


```
plot(dat$age2, dat$cigage)
```



```
cor(dat$age2, dat$iralcage)
```

```
## [1] 0.07253557
```

```
plot(dat$age2, dat$iralcage)
```




People in the survey tended to use alcohol earlier than other substances. This can be seen by the scatter plots. Additionally, alcohol has the lowest correlation coefficient.

Problem 5: Sexual attraction

What does the distribution of sexual attraction look like? Is this what you expected?

```
counts <- table(dat$sexattract)
barplot(counts, main= "Sexual Attraction Preferences Based on Small Sample from NSDUH 2019", xlab= "Cat
```

Sexual Attraction Preferences Based on Small Sample from NSDUH 2




According to this bargraph, most people answered 1, which indicates preference for the opposite sex. This does not surprise me, as most people in the United States are heterosexual, however I would expect more people to answer 3 or 5.

What is the distribution of sexual attraction by gender?

```
counts <- table(dat$irsex, dat$sexattract)
barplot(counts, main= "Sexual Attraction Preferences Seperated by Gender ", xlab="Categories", ylab= "N
```

Sexual Attraction Preferences Seperated by Gender




Group 1, which is heterosexual, is mainly composed of males. While group 3, which is bisexual is completely female. Group 6, which is homosexual, is also completely male.

Problem 6: English speaking

What does the distribution of English speaking look like in the sample? Is this what you might expect for a random sample of the US population?

```
counts <- table(dat$speakengl)
barplot(counts, main= "How Well Sample of Respondents of NSDUH 2019 Speak English", xlab="Categories", ylab="Number of Respondents")
```

How Well Sample of Respondents of NSDUH 2019 Speak English




Most participants speak English very well, very few speak well and almost none speak not well or none at all. If this survey was only given out in English, this distribution makes sense. But in the US population, while the majority of people speak English, many do not speak it “very well.”

Are there more English speaker female or males?

```
counts <- table(dat$irsex, dat$peakengl)
barplot(counts, main= "How Well Sample of Respondents of NSDUH 2019 Speak English Seperated by Gender",
```

Well Sample of Respondents of NSDUH 2019 Speak English Seperated



There are a similar number of males and females who speak English “very well” There are many more females who speak English “Well,” so overall there are more females then males who speak English.

Exam 1

Instructions

- Create a folder in your computer (a good place would be under Crim 250, Exams).
- Download the dataset from the Canvas website (fatal-police-shootings-data.csv) onto that folder, and save your Exam 1.Rmd file in the same folder.
- Download the README.md file. This is the codebook.
- Load the data into an R data frame.

```
dat <- read.csv(file = 'fatal-police-shootings-data.csv')
head(dat)
```

##	id	name	date	manner_of_death	armed	age	gender	race
## 1	3	Tim Elliot	2015-01-02	shot	gun	53	M	A
## 2	4	Lewis Lee Lembke	2015-01-02	shot	gun	47	M	W
## 3	5	John Paul Quintero	2015-01-03	shot and Tasered	unarmed	23	M	H
## 4	8	Matthew Hoffman	2015-01-04	shot toy weapon		32	M	W

```
## 5 9 Michael Rodriguez 2015-01-04 shot nail gun 39 M H
## 6 11 Kenneth Joe Brown 2015-01-04 shot gun 18 M W
##      city state signs_of_mental_illness threat_level flee
## 1      Shelton WA True attack Not fleeing
## 2      Aloha OR False attack Not fleeing
## 3      Wichita KS False other Not fleeing
## 4 San Francisco CA True attack Not fleeing
## 5      Evans CO False attack Not fleeing
## 6      Guthrie OK False attack Not fleeing
## body_camera longitude latitude is_geocoding_exact
## 1      False -123.122 47.247 True
## 2      False -122.892 45.487 True
## 3      False -97.281 37.695 True
## 4      False -122.422 37.763 True
## 5      False -104.692 40.384 True
## 6      False -97.423 35.877 True
```

Problem 1 (10 points)

a. Describe the dataset.

This is the source: <https://github.com/washingtonpost/data-police-shootings>. Write two sentences (max.) about this.

This data set is data collected by the Washington Post. It gives information about how many fatal shootings there were by an active duty police officer, and gives details about the victim (example: armed or unarmed.)

b. How many observations are there in the data frame?

```
nrow(dat)
```

```
## [1] 6594
```

There are 6594 observations in the data frame.

c. Look at the names of the variables in the data frame. Describe what “body_camera”, “flee”, and “armed” represent, according to the codebook. Again, only write one sentence (max) per variable.

```
names(dat)
```

```
## [1] "id" "name"
## [3] "date" "manner_of_death"
## [5] "armed" "age"
## [7] "gender" "race"
## [9] "city" "state"
## [11] "signs_of_mental_illness" "threat_level"
## [13] "flee" "body_camera"
## [15] "longitude" "latitude"
## [17] "is_geocoding_exact"
```


“Body_camera” indicates whether or not the officer was wearing a body camera that may have recorded part of the incident. “flee” represents whether or not the victim was moving away from the officers and how they fled. “armed” represents whether or not the victim was armed with something the officers believed could inflict harm and specifies what, if anything, the victims were carrying.

d. What are three weapons that you are surprised to find in the “armed” variable? Make a table of the values in “armed” to see the options.

```
table(dat$armed)
```

```
##
##
##          air conditioner
##          207            1
##          air pistol      Airsoft pistol
##          1              3
##          ax              barstool
##          24             1
##          baseball bat    baseball bat and bottle
##          20             1
## baseball bat and fireplace poker    baseball bat and knife
##          1              1
##          baton           BB gun
##          6              15
##          BB gun and vehicle    bean-bag gun
##          1              1
##          beer bottle          binoculars
##          3              1
##          blunt object         bottle
##          5              1
##          bow and arrow        box cutter
##          1              13
##          brick               car, knife and mace
##          2              1
##          carjack             chain
##          1              3
##          chain saw           chainsaw
##          2              1
##          chair              claimed to be armed
##          4              1
##          contractor's level    cordless drill
##          1              1
##          crossbow             crowbar
##          9              5
##          fireworks            flagpole
##          1              1
##          flashlight           garden tool
##          2              2
##          glass shard          grenade
##          4              1
##          gun                 gun and car
##          3798              12
##          gun and knife       gun and machete
```

##	22	3
##	gun and sword	gun and vehicle
##	1	17
##	guns and explosives	hammer
##	3	18
##	hand torch	hatchet
##	1	14
##	hatchet and gun	ice pick
##	2	1
##	incendiary device	knife
##	2	955
##	knife and vehicle	lawn mower blade
##	1	2
##	machete	machete and gun
##	51	1
##	meat cleaver	metal hand tool
##	6	2
##	metal object	metal pipe
##	5	16
##	metal pole	metal rake
##	4	1
##	metal stick	microphone
##	3	1
##	motorcycle	nail gun
##	1	1
##	oar	pellet gun
##	1	3
##	pen	pepper spray
##	1	2
##	pick-axe	piece of wood
##	4	7
##	pipe	pitchfork
##	7	2
##	pole	pole and knife
##	3	2
##	railroad spikes	rock
##	1	7
##	samurai sword	scissors
##	4	9
##	screwdriver	sharp object
##	16	14
##	shovel	spear
##	7	2
##	stapler	straight edge razor
##	1	5
##	sword	Taser
##	23	34
##	tire iron	toy weapon
##	4	226
##	unarmed	undetermined
##	421	188
##	unknown weapon	vehicle
##	82	213
##	vehicle and gun	vehicle and machete

```
##                8                1
##      walking stick      wasp spray
##                1                1
##      wrench
##                1
```


I was surprised to find binoculars, air conditioners and a motorcycle, because these are not things that I would typically think of as weapons. I wonder who determined that these items were a threat.

Problem 2 (10 points)

a. Describe the age distribution of the sample. Is this what you would expect to see?

```
hist(dat$age, main= "Frequency of ages of police shooting victims from Washington Post Data", xlab = "Ages of victims", ylab = "Frequency")
```

Frequency of ages of police shooting victims from Washington Post Data



The age distribution is right skewed. This means that most of the victims were younger (20-40) This is what I would expect to see, because police shootings typically happen to people who police “perceive as threats” and very old people typically won’t be seen as threats.

b. To understand the center of the age distribution, would you use a mean or a median, and why? Find the one you picked.

```
counts <- table(dat$age)

median(counts)
```

```
## [1] 66.5
```

Because this data is right skewed I used a median to find the center of age distribution. Outliers can have more of an impact on the mean, which would not make it as useful as a statistician.

c. Describe the gender distribution of the sample. Do you find this surprising?

```
table(dat$gender)
```

```
##
##      F      M
##  3  293 6298
```

The gender distribution is not even. There are far more men than women. This makes sense because police usually perceive men as more of a threat than women. Additionally, men are arrested more than women overall according to the FBI.

Problem 3 (10 points)

a. How many police officers had a body camera, according to news reports? What proportion is this of all the incidents in the data? Are you surprised that it is so high or low?

```
table(dat$body_camera)
```

```
##
## False  True
##  5684   910
```

```
sum(table(dat$body_camera))
```

```
## [1] 6594
```

```
5684/6594
```

```
## [1] 0.8619958
```

There are many more people who were not wearing body cams than people who were. There were 5684 people not wearing body cams out of the 6594 recorded, which means 86.2 percent of officers did not have their cameras on. I would think that there would be more people wearing body cams in data recorded since 2015, because there has been a big push for more body cams in the recent years. However, many police officers do not turn on their body cameras.

b. In how many of the incidents was the victim fleeing? What proportion is this of the total number of incidents in the data? Is this what you would expect?

```
table(dat$flee)
```

```
##
##           Car      Foot Not fleeing      Other
##           491     1058         845       3952        248
```

```
1058 + 845 + 3952 + 248
```

```
## [1] 6103
```

```
(1058 + 845 + 248)/6103
```

```
## [1] 0.3524496
```

There are 491 pieces of missing data I excluded these from my calculations. I assume this means that the people who collected the data do not know whether the victims fled, but they did not make it clear. Additionally, I am assuming that “other” means people who fled in a different way than car or foot. There are a total of 6103 pieces of data (minus the data that was missing) and there are 845 people who fled on foot, 1058 by car and 248 by other. This means that only 35 percent of victims attempted to flee. I would expect it to be more, because police officers should not have to shoot people who are not fleeing. At the same time, if a victim is fleeing they may not be an imminent danger to a police officer (because they are running, not attacking), so it is possible the officers wouldn’t shoot.

Problem 4 (10 points) - Answer only one of these (a or b).

- a. Describe the relationship between the variables “body camera” and “flee” using a stacked barplot. What can you conclude from this relationship?

Hint 1: The categories along the x-axis are the options for “flee”, each bar contains information about whether the police officer had a body camera (vertically), and the height along the y-axis shows the frequency of that category).

Hint 2: Also, if you are unsure about the syntax for barplot, run ?barplot in R and see some examples at the bottom of the documentation. This is usually a good way to look up the syntax of R code. You can also Google it.

Your answer here.

- b. Describe the relationship between age and race by using a boxplot. What can you conclude from this relationship?


Hint 1: The categories along the x-axis are the race categories and the height along the y-axis is age.

Hint 2: Also, if you are unsure about the syntax for boxplot, run ?boxplot in R and see some examples at the bottom of the documentation. This is usually a good way to look up the syntax of R code. You can also Google it.

```
table(dat$race)
```

```
##
##           A      B      H      N      O      W
## 752  106 1553 1083    91    47 2962
```

```
plot(factor(dat$race), dat$age, main= "Barplot of Age of Victims and Race", xlab= "Race in Catagories", ylab= "Age of victims", names=c("unkown", "Asian", "Black", "Hispanic", "NA ", "Other", "White"))
```



NOTE there is missing data. The first bar plot represents “unknown” data. Black= Black Non-Hispanic, White= White, non-Hispanic, NA= Native American

There is a clear distinction between the median age victims who are white and every other race (excluding unknown). The median age of victims who are white are much higher than the median age of victims of the other races. Additionally, the ages of white victims are much more evenly dispersed. The box is close to the middle which shows a more even dispersion, and the box is longer which shows more dispersed data within Q2 and Q3. Native Americans and black victims have the lowest median age, although there are many outliers who are black. This data shows that victims or color tend to be younger than white victims of police shootings.

Extra credit (10 points)

a. What does this code tell us?

```
mydates <- as.Date(dat$date)
head(mydates)
(mydates[length(mydates)] - mydates[1])
```

2458/365

This code tells us the number of days of data we have. The data has been being kept since 2015, and it is now 2021. So it makes sense that the data is of 2458 days (6.7 years)

b. On Friday, a new report was published that was described as follows by The Guardian: “More than half of US police killings are mislabelled or not reported, study finds.” Without reading this article now (due to limited time), why do you think police killings might be mislabelled or underreported?

It seems that police killings are underreported in federal data bases. According to the Washington Post Study, the FBI and CDC does not log as many fatal shootings as there actually are. Police killings are most likely under reported because it is a liability to police forces when one of their officers shoots and kills a victim. Additionally, as this data shows, minority victims tend to be younger, which demonstrates possible racial prejudice. Police killings are most likely misreported because people may blame death on other factors than a gun wound. If the cause of death is reported differently, then a victim would not be counted towards a police shooting death.

c. Regarding missing values in problem 4, do you see any? If so, do you think that’s all that’s missing from the data?

In problem 4 the missing value is an unknown race. This differs from the race category of ‘other,’ which shows victims who do not identify as one of the specific racial categories laid out by the Washington Post. I assume that there is much more information missing in general. There is missing data about whether victims fled. There are probably many missing values of signs of mental illness, because that is not easy data to collect. Additionally, the data that the Washington Post collected is from many different sources, none of which have all the pieces data they wish to collect. Overall, there is probably a lot of data and variables missing that are important to understanding the full scope of the circumstances of the shooting. This data just shows us a very small lens of information.

Assignment 3

title: “Assignment 3” author: “Niko Amber” date: “Today”’s date here: 10/28/2021’ output: pdf_document:
default html_document: default —

I worked with Bella Werneck

This assignment is due on Canvas on Wednesday 10/27/2021 before class, at 10:15 am. Include the name of anyone with whom you collaborated at the top of the assignment.

Submit your responses as either an HTML file or a PDF file on Canvas. Also, please upload it to your website.

Save the file (found on Canvas) crime_simple.txt to the same folder as this file (your Rmd file for Assignment 3).

Load the data.

```
library(readr)
library(knitr)
dat_crime <- read_delim("/Users/nikoamber/Library/Mobile Documents/com~apple~CloudDocs/Penn/Freshman Fa
```

```
## Rows: 47 Columns: 14
```

```
## -- Column specification -----
```

```
## Delimiter: "\t"
```

```
## dbl (14): R, Age, S, Ed, Ex0, Ex1, LF, M, N, NW, U1, U2, W, X
```

```
##
```

```
## i Use 'spec()' to retrieve the full column specification for this data.
```

```
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

This is a dataset from a textbook by Brian S. Everitt about crime in the US in 1960. The data originate from the Uniform Crime Report of the FBI and other government sources. The data for 47 states of the USA are given.

Here is the codebook:

R: Crime rate: # of offenses reported to police per million population

Age: The number of males of age 14-24 per 1000 population

S: Indicator variable for Southern states (0 = No, 1 = Yes)

Ed: Mean of years of schooling x 10 for persons of age 25 or older

Ex0: 1960 per capita expenditure on police by state and local government

Ex1: 1959 per capita expenditure on police by state and local government

LF: Labor force participation rate per 1000 civilian urban males age 14-24

M: The number of males per 1000 females

N: State population size in hundred thousands

NW: The number of non-whites per 1000 population

U1: Unemployment rate of urban males per 1000 of age 14-24

U2: Unemployment rate of urban males per 1000 of age 35-39

W: Median value of transferable goods and assets or family income in tens of \$

X: The number of families per 1000 earning below 1/2 the median income

We are interested in checking whether the reported crime rate (# of offenses reported to police per million population) and the average education (mean number of years of schooling for persons of age 25 or older) are related.

1. How many observations are there in the dataset? To what does each observation correspond?

```
nrow(dat_crime)
```

```
## [1] 47
```



```
ncol(dat_crime)
```

```
## [1] 14
```

```
summary(dat_crime)
```

```
##           R           Age           S           Ed
## Min.      : 34.20   Min.    :119.0   Min.      :0.0000   Min.      : 87.0
## 1st Qu.: 65.85   1st Qu.:130.0   1st Qu.:0.0000   1st Qu.: 97.5
## Median : 83.10   Median :136.0   Median :0.0000   Median :108.0
## Mean      : 90.51   Mean     :138.6   Mean      :0.3404   Mean      :105.6
## 3rd Qu.:105.75   3rd Qu.:146.0   3rd Qu.:1.0000   3rd Qu.:114.5
## Max.      :199.30   Max.      :177.0   Max.      :1.0000   Max.      :122.0
##           Ex0          Ex1          LF          M
## Min.      : 45.0   Min.      : 41.00   Min.      :480.0   Min.      : 934.0
## 1st Qu.: 62.5   1st Qu.: 58.50   1st Qu.:530.5   1st Qu.: 964.5
## Median : 78.0   Median : 73.00   Median :560.0   Median : 977.0
## Mean      : 85.0   Mean      : 80.23   Mean      :561.2   Mean      : 983.0
## 3rd Qu.:104.5   3rd Qu.: 97.00   3rd Qu.:593.0   3rd Qu.: 992.0
## Max.      :166.0   Max.      :157.00   Max.      :641.0   Max.      :1071.0
##           N           NW           U1           U2
## Min.      : 3.00   Min.      : 2.0   Min.      : 70.00   Min.      :20.00
## 1st Qu.: 10.00   1st Qu.: 24.0   1st Qu.: 80.50   1st Qu.:27.50
## Median : 25.00   Median : 76.0   Median : 92.00   Median :34.00
## Mean      : 36.62   Mean      :101.1   Mean      : 95.47   Mean      :33.98
## 3rd Qu.: 41.50   3rd Qu.:132.5   3rd Qu.:104.00   3rd Qu.:38.50
## Max.      :168.00   Max.      :423.0   Max.      :142.00   Max.      :58.00
##           W           X
## Min.      :288.0   Min.      :126.0
## 1st Qu.:459.5   1st Qu.:165.5
## Median :537.0   Median :176.0
## Mean      :525.4   Mean      :194.0
## 3rd Qu.:591.5   3rd Qu.:227.5
## Max.      :689.0   Max.      :276.0
```

```
head(dat_crime)
```

```
## # A tibble: 6 x 14
##       R   Age   S   Ed  Ex0  Ex1  LF   M   N   NW  U1  U2   W
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  79.1  151     1    91    58    56   510   950    33   301   108    41   394
## 2 164.   143     0   113   103    95   583  1012    13   102    96    36   557
## 3  57.8  142     1    89    45    44   533   969    18   219    94    33   318
## 4 197.   136     0   121   149   141   577   994   157    80   102    39   673
## 5 123.   141     0   121   109   101   591   985    18    30    91    20   578
## 6  68.2  121     0   110   118   115   547   964    25    44    84    29   689
## # ... with 1 more variable: X <dbl>
```


There are 47 observations in the data frame. Each row gives the data for each of the 47 states included in this data set. Each of the 47 rows includes data from each state which gives information about the 14 variables explained in the code book.

out 4 variables of data from each state including, level of education, age and whether or not the states are in the South.

2. Draw a scatterplot of the two variables. Calculate the correlation between the two variables. Can you come up with an explanation for this relationship?

```
plot(dat_crime$Ed, dat_crime$R, xlab= "Average Years of Schooling x 10 for people over 25", ylab= "Crime Rate per state According to 1960 UCR Data")
```

Relationship between Level of Education and Crime Rate for 47 states



```
head(dat_crime)
```

```
## # A tibble: 6 x 14
##       R   Age    S    Ed  Ex0  Ex1   LF    M    N   NW   U1   U2    W
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  79.1  151     1    91    58    56   510   950    33   301   108    41   394
## 2 164.   143     0   113   103    95   583  1012    13   102    96    36   557
## 3  57.8  142     1    89    45    44   533   969    18   219    94    33   318
## 4 197.   136     0   121   149   141   577   994   157    80   102    39   673
## 5 123.   141     0   121   109   101   591   985    18    30    91    20   578
## 6  68.2  121     0   110   118   115   547   964    25    44    84    29   689
## # ... with 1 more variable: X <dbl>
```

```
cor(dat_crime$Ed, dat_crime$R)
```

```
## [1] 0.3228349
```

The correlation efficient is .3228349 which indicates a slight positive relationship between crime rate and level of education. This is not what I would have expected. It is important to note this is correlation not

causation, so higher levels of education do not result in more crime. This data includes all crime types. It is possible that there are a lot of white collar criminals which typically have higher levels of education. Additionally, this data is based on the crime reported, not all crime that has happened which could have an impact on the correlation. The correlation is not strong, but I would like to know why there is a positive correlation at all.

3. Regress reported crime rate (y) on average education (x) and call this linear model `crime.lm` and write the summary of the regression by using this code, which makes it look a little nicer `{r, eval=FALSE} summary(crime.lm)$coef, digits = 2`.

```
lm( formula = dat_crime$R ~ dat_crime$Ed, data = (dat_crime) )
```

```
##
## Call:
## lm(formula = dat_crime$R ~ dat_crime$Ed, data = (dat_crime))
##
## Coefficients:
## (Intercept)  dat_crime$Ed
##      -27.397      1.116
```

```
lmCrime <- lm( formula = dat_crime$R ~ dat_crime$Ed, data = (dat_crime) )
summary(lmCrime)
```

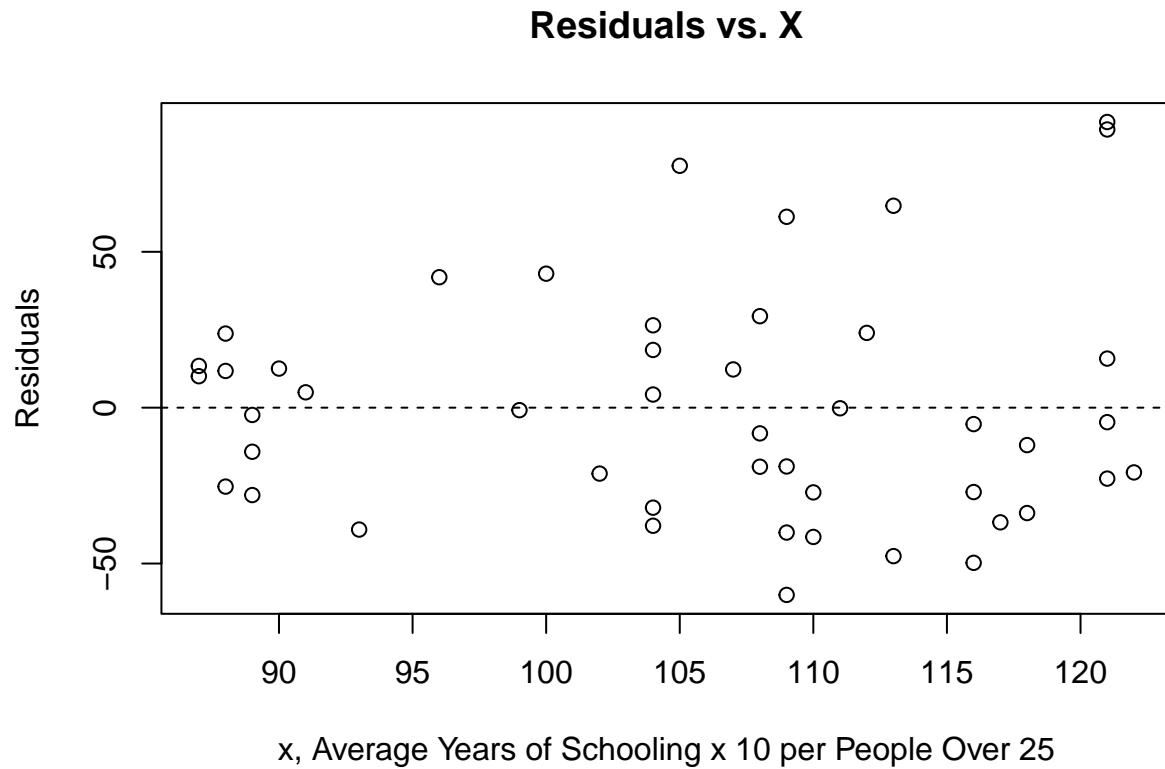
```
##
## Call:
## lm(formula = dat_crime$R ~ dat_crime$Ed, data = (dat_crime))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -60.061 -27.125  -4.654   17.133   91.646
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -27.3967     51.8104  -0.529   0.5996
## dat_crime$Ed   1.1161     0.4878   2.288   0.0269 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 37.01 on 45 degrees of freedom
## Multiple R-squared:  0.1042, Adjusted R-squared:  0.08432
## F-statistic: 5.236 on 1 and 45 DF,  p-value: 0.02688
```

```
coef(lmCrime, digit = 2)
```

```
## (Intercept) dat_crime$Ed
##      -27.396654      1.116121
```


4. Are the four assumptions of linear regression satisfied? To answer this, draw the relevant plots. (Write a maximum of one sentence per assumption.)

```
plot(dat_crime$Ed, lmCrime$residuals, main= "Residuals vs. X", xlab = "x, Average Years of Schooling x 10 per People Over 25", ylab = "Residuals", abline(h = 0, lty = "dashed"))
```




The linearity assumption is satisfied because the residuals vs X plot does not have a pattern around the horizontal line $y=0$.

```
plot(lmCrime, which = 1)
```




The independent assumption is satisfied because not only does the residuals vs X plot not have a pattern around the horizontal line x, also the residuals vs fitted shows no pattern.

```
plot(lmCrime, which = 3)
```



The equal variance assumption helps us analyze the homoskedasticity assumption, which is satisfied because the red line is close to being horizontal.

```
plot(lmCrime, which=2)
```



The normal QQ plot tells us that the normal population is slightly right skewed, there is a light tail, which means the distribution isn't normal.

5. Is the relationship between reported crime and average education statistically significant? Report the estimated coefficient of the slope, the standard error, and the p-value. What does it mean for the relationship to be statistically significant?

```
summary(lmCrime)
```

```
##
## Call:
## lm(formula = dat_crime$R ~ dat_crime$Ed, data = (dat_crime))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -60.061 -27.125  -4.654  17.133  91.646
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -27.3967    51.8104  -0.529   0.5996
## dat_crime$Ed   1.1161     0.4878   2.288   0.0269 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 37.01 on 45 degrees of freedom
## Multiple R-squared:  0.1042, Adjusted R-squared:  0.08432
## F-statistic: 5.236 on 1 and 45 DF, p-value: 0.02688
```

The p value is less than 0.05 so the relationship between reported crime and average education is statistically significant. A relationship is statistically significant when you can reject the null hypothesis.

6. How are reported crime and average education related? In other words, for every unit increase in average education, how does reported crime rate change (per million) per state?

For every unit increase in average education, the crime rate changes 1.1161 per million per state.

7. Can you conclude that if individuals were to receive more education, then crime will be reported more often? Why or why not?

No. Correlation does not mean causation. Just because there seems to be a positive correlation between level of education and reported crime rates, does not mean that more education CAUSES more crime.